

CS 285 Assignment 2: Policy Gradients Report

1 Experiment 1: CartPole (Section 3.2)

1.1 Learning Curves

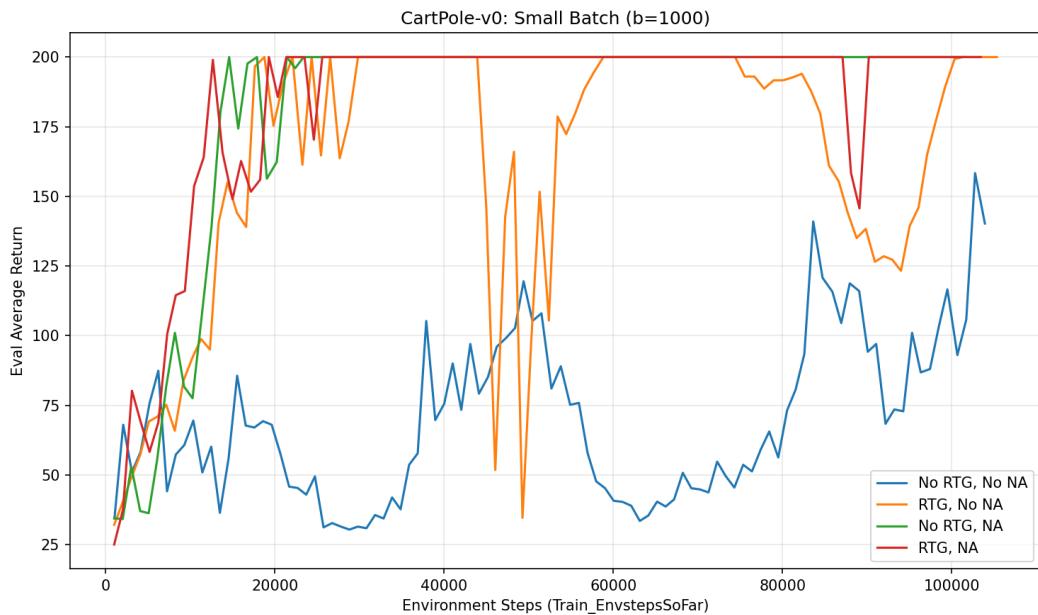


Figure 1: CartPole-v0 learning curves with small batch size ($b = 1000$). The x -axis is the number of environment steps (Train_EnvstepsSoFar).

All configurations in both large and small batch settings converge to the maximum return of 200, except for the vanilla configurations without reward-to-go and without advantage normalization, which show the worst performance: in the small batch setting, vanilla (No RTG, No NA) only reaches ~ 140 at the end, and in the large batch setting, vanilla (No RTG, No NA) only reaches ~ 102 .

1.2 Analysis Questions

Which value estimator has better performance without advantage normalization: the trajectory-centric one, or the one using reward-to-go?

The reward-to-go (RTG) estimator has better performance without advantage normalization. In the small batch setting, RTG reaches 200 while the trajectory-centric estimator only reaches

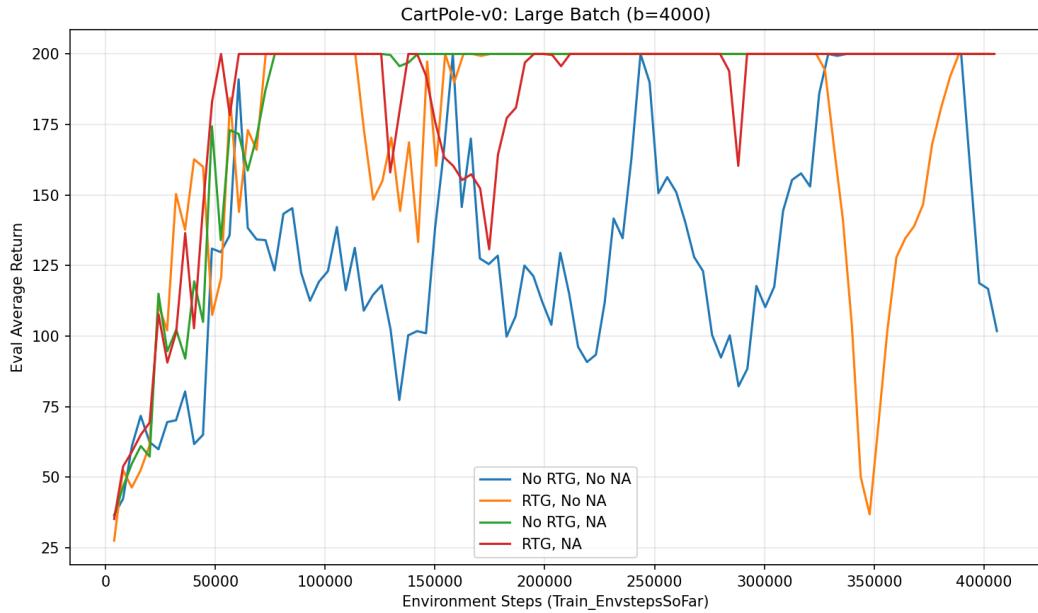


Figure 2: CartPole-v0 learning curves with large batch size ($b = 4000$). The x -axis is the number of environment steps (Train_EnvstepsSoFar).

~ 140 . The same pattern holds in the large batch setting. This is expected because RTG reduces variance by removing rewards from past timesteps that cannot be affected by the current action.

Between the two value estimators, why do you think one is generally preferred over the other?

Reward-to-go is generally preferred because it exploits the causality structure of the MDP: actions at time t can only affect rewards at times $t' \geq t$. By excluding past rewards from the return estimate, RTG reduces variance without introducing any bias. The trajectory-centric estimator includes rewards from the past which act purely as noise in the gradient estimate.

Did advantage normalization help?

Yes. Advantage normalization consistently improves or matches performance across all settings. In particular, it helps the trajectory-centric estimator (No RTG) converge much more reliably to the maximum return of 200. By normalizing to zero mean and unit variance, the gradient updates become more stable regardless of the scale of the returns.

Did the batch size make an impact?

Yes, increasing the batch size from 1000 to 4000 generally improves learning stability and convergence speed in terms of number of iterations. With a larger batch, the gradient estimate has lower variance. However, each iteration also collects more data, so the total number of environment steps is higher. The large batch is most helpful for the less effective configurations (e.g., No RTG, No NA), making them more stable.

1.3 Command Line Configurations

```
# Small batch experiments (b=1000)
uv run src/scripts/run.py --env_name CartPole-v0 -n 100 -b 1000 --exp_name
    cartpole
```

```

uv run src/scripts/run.py --env_name CartPole-v0 -n 100 -b 1000 -rtg --
    exp_name cartpole_rtg
uv run src/scripts/run.py --env_name CartPole-v0 -n 100 -b 1000 -na --
    exp_name cartpole_na
uv run src/scripts/run.py --env_name CartPole-v0 -n 100 -b 1000 -rtg -na
    --exp_name cartpole_rtg_na

# Large batch experiments (b=4000)
uv run src/scripts/run.py --env_name CartPole-v0 -n 100 -b 4000 --exp_name
    cartpole_lb
uv run src/scripts/run.py --env_name CartPole-v0 -n 100 -b 4000 -rtg --
    exp_name cartpole_lb_rtg
uv run src/scripts/run.py --env_name CartPole-v0 -n 100 -b 4000 -na --
    exp_name cartpole_lb_na
uv run src/scripts/run.py --env_name CartPole-v0 -n 100 -b 4000 -rtg -na
    --exp_name cartpole_lb_rtg_na

```

2 Experiment 2: HalfCheetah (Section 4.2)

2.1 Learning Curves

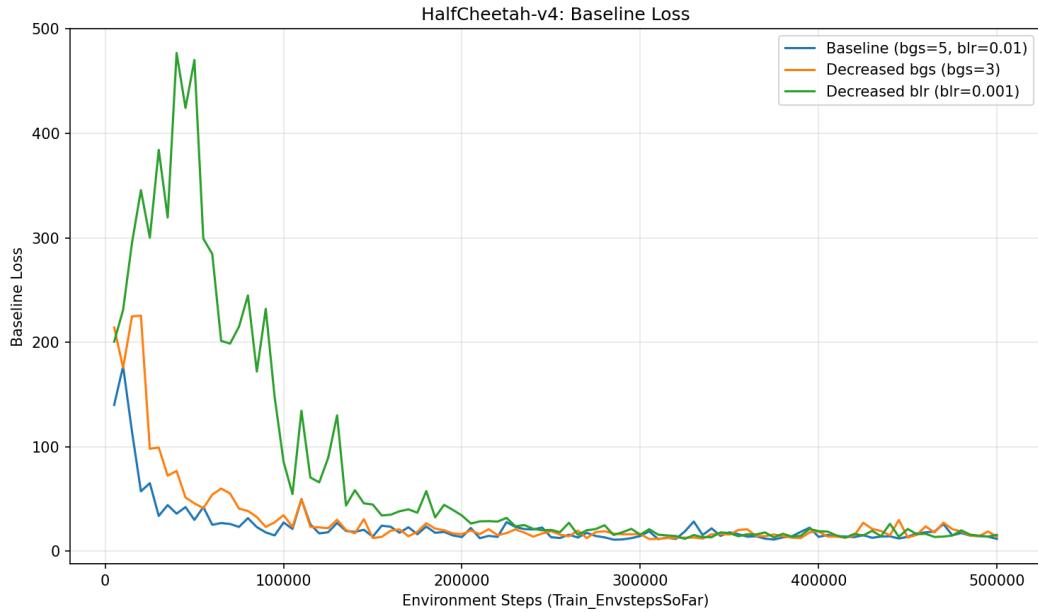


Figure 3: HalfCheetah-v4: Baseline loss curves comparing the default baseline configuration with decreased baseline gradient steps ($bgs=3$) and decreased baseline learning rate ($blr=0.001$).

The baselined policy gradient achieves a final eval return of ~ 444 , well above the threshold of 300. The no-baseline version finishes at ~ 196 , demonstrating the significant benefit of using a learned value function baseline for variance reduction.

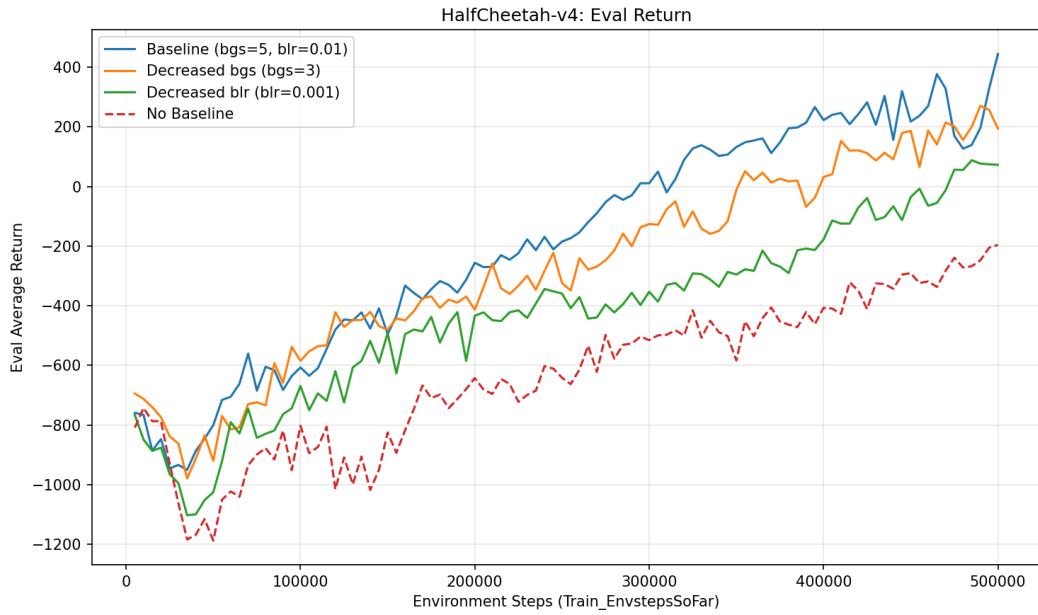


Figure 4: HalfCheetah-v4: Eval return curves. The baselined version ($bgs=5$, $blr=0.01$) achieves an average return above 300 at the end of training, while the no-baseline version performs significantly worse.

2.2 Effect of Decreased Baseline Gradient Steps / Learning Rate

We ran two additional experiments varying the baseline hyperparameters:

- **Decreased bgs (bgs=3 vs. default 5):** The baseline loss curve is slightly higher (worse fit), and the final eval return drops to ~ 194 . With fewer gradient steps, the baseline network cannot fit the value function as accurately, leading to a less effective baseline and noisier policy gradients.
- **Decreased blr (blr=0.001 vs. default 0.01):** The baseline loss remains high throughout training, and the final eval return drops to ~ 73 . The smaller learning rate severely limits the baseline network’s ability to learn, resulting in poor variance reduction and degraded policy performance.

Both modifications hurt performance, confirming that the quality of the baseline fit directly affects policy learning.

2.3 Command Line Configurations

```
# No baseline
uv run src/scripts/run.py --env_name HalfCheetah-v4 -n 100 -b 5000 -eb
3000 \
-rtg --discount 0.95 -lr 0.01 --exp_name cheetah

# Baseline (default)
uv run src/scripts/run.py --env_name HalfCheetah-v4 -n 100 -b 5000 -eb
3000 \
```

```

-rtg --discount 0.95 -lr 0.01 --use_baseline -blr 0.01 -bgs 5 \
--exp_name cheetah_baseline

# Decreased baseline gradient steps (bgs=3)
uv run src/scripts/run.py --env_name HalfCheetah-v4 -n 100 -b 5000 -eb
3000 \
-rtg --discount 0.95 -lr 0.01 --use_baseline -blr 0.01 -bgs 3 \
--exp_name cheetah_baseline_decreased_bgs

# Decreased baseline learning rate (blr=0.001)
uv run src/scripts/run.py --env_name HalfCheetah-v4 -n 100 -b 5000 -eb
3000 \
-rtg --discount 0.95 -lr 0.01 --use_baseline -blr 0.001 -bgs 5 \
--exp_name cheetah_baseline_decreased_bgr

```

3 Experiment 3: LunarLander (Section 5)

3.1 Learning Curves

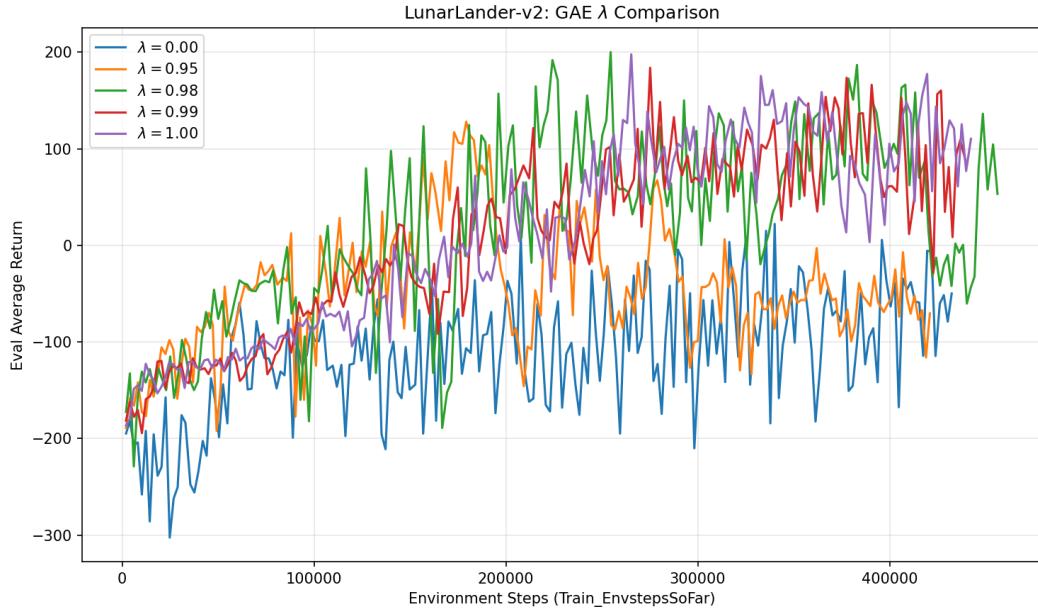


Figure 5: LunarLander-v2: Learning curves for GAE- λ with $\lambda \in \{0, 0.95, 0.98, 0.99, 1\}$.

3.2 Analysis

How did λ affect task performance?

$\lambda = 0$ performs the worst, with max eval return of only ~ 22 . This is because $\lambda = 0$ corresponds to using only the 1-step TD error δ_t as the advantage estimate, which has high bias. $\lambda = 0.95$ improves significantly (max ~ 128) but still does not reach 150. The best performance is achieved by $\lambda = 0.98$ (max ~ 200) and $\lambda = 1.0$ (max ~ 198), with $\lambda = 0.99$ also performing well (max ~ 184).

Overall, higher values of λ lead to better performance on this task, suggesting that variance is less of a concern than bias for LunarLander-v2.

What does $\lambda = 0$ correspond to? What about $\lambda = 1$?

$\lambda = 0$ corresponds to using only the 1-step TD error $\delta_t = r_t + \gamma V(s_{t+1}) - V(s_t)$ as the advantage estimate, which is low variance but high bias. $\lambda = 1$ corresponds to using the full Monte Carlo return minus the baseline (i.e., $Q^\pi(s_t, a_t) - V^\pi(s_t)$), which is unbiased but higher variance. In LunarLander-v2, the high bias of $\lambda = 0$ hurts significantly more than the higher variance of $\lambda = 1$, so larger λ values are preferred.

3.3 Command Line Configurations

```
uv run src/scripts/run.py --env_name LunarLander-v2 --ep_len 1000 \
--discount 0.99 -n 200 -b 2000 -eb 2000 -l 3 -s 128 -lr 0.001 \
--use_reward_to_go --use_baseline --gae_lambda <lambda> \
--exp_name lunar_lander_lambda<lambda>
# where <lambda> in {0, 0.95, 0.98, 0.99, 1}
```

4 Experiment 4: InvertedPendulum (Section 6)

4.1 Learning Curves




Figure 6: InvertedPendulum-v4: Learning curves for different hyperparameter configurations. The best configuration reaches a return of 1000 within $\sim 35K$ environment steps.

4.2 Best Hyperparameters

The best configuration reaches the maximum return of 1000 within $\sim 35,000$ environment steps (well under the 100K budget). The key hyperparameters are:

- **Learning rate ($lr = 0.02$):** A higher policy learning rate was the most important factor, allowing faster convergence.
- **Batch size ($b = 1000$):** A smaller batch size (compared to the default 5000) means each policy gradient iteration uses fewer environment steps, enabling more frequent policy updates within the step budget.
- **Reward-to-go, baseline, GAE ($\lambda = 0.98$), and advantage normalization:** All variance reduction techniques were enabled, providing stable and efficient gradient estimates.
- **Discount ($\gamma = 0.99$):** A high discount factor is important for this continuing task.
- **Network size (layer_size=128):** A larger network than the default 64 was used.

The most impactful hyperparameters were the **learning rate** and **batch size**. Increasing the learning rate from 0.001 to 0.02 dramatically sped up convergence, and using a smaller batch size allowed more policy updates per environment step.

4.3 Command Line Configuration

```
# Best configuration (reaches 1000 at ~35K steps)
uv run src/scripts/run.py --env_name InvertedPendulum-v4 -n 200 -b 1000 \
-eb 800 -rtg --use_baseline -blr 0.005 -bgs 5 --gae_lambda 0.98 \
-na --discount 0.99 -lr 0.02 -s 128 --exp_name pendulum
```